

# On Graphs, GPUs, and Blind Dating

A Workload to Processor Matchmaking Quest

Abdullah Gharaibeh, Lauro Beltro Costa, Elizeu Santos-Neto, Matei  
Ripeanu

Department of Electrical and Computer Engineering  
The University of British Columbia

presented by Ulla Aeschbacher

Institute of Computer Science, ETH Zürich

17.4.18

# Goal of Paper

- ▶ Explore effectiveness of graph partitioning strategies and workload allocation schemes to maximize performance
- ▶ Give evaluation of said performance

# Overview

1. Goal of Paper
2. Related Work and Background
3. Characteristics
4. Partitioning Strategy
5. Evaluation
6. Summary

## Related Work

- ▶ September 2012: A Yoke of Oxen and a Thousand Chickens for Heavy Lifting Graph Processing

# Related Work

- ▶ September 2012: A Yoke of Oxen and a Thousand Chickens for Heavy Lifting Graph Processing
- ▶ May 2013: On Graphs, GPUs, and Blind Dating

# Related Work

- ▶ September 2012: A Yoke of Oxen and a Thousand Chickens for Heavy Lifting Graph Processing
- ▶ May 2013: On Graphs, GPUs, and Blind Dating
- ▶ November 2013: The Energy Case for Graph Processing on Hybrid CPU and GPU Systems

# TOTEM Computation Model

- ▶ Processing divided into rounds
- ▶ Rounds consisting of three phases
  - ▶ Computation
  - ▶ Communication
  - ▶ Synchronization

# Characteristics of (Most) Real-World Graph Workloads

- ▶ Modest processing per vertex per round



# Characteristics of (Most) Real-World Graph Workloads

- ▶ Modest processing per vertex per round
- ▶ Poor locality

# Characteristics of (Most) Real-World Graph Workloads

- ▶ Modest processing per vertex per round
- ▶ Poor locality
- ▶ Power-law degree distribution

# Characteristics of Hybrid System

## CPU Intel Nehalem Xeon X5650

- ▶  $2 \times$  core frequency
  - ▶  $24.6 \times$  more memory
  - ▶  $6 \times$  more cache
  - ▶ retailed at 999\$
- Can process large graphs

## GPU Nvidia Tesla C2075

- ▶  $16 \times$  more hardware threads
  - ▶  $4.5 \times$  faster memory access
  - ▶ retailed at 599\$
- Can process graphs in parallel

# Required Features for Graph Partitioning Strategies

- ▶ Low space and time complexity

# Required Features for Graph Partitioning Strategies

- ▶ Low space and time complexity
- ▶ Handles scale-free graphs

# Required Features for Graph Partitioning Strategies

- ▶ Low space and time complexity
- ▶ Handles scale-free graphs
- ▶ Handles large graphs

# Required Features for Graph Partitioning Strategies

- ▶ Low space and time complexity
- ▶ Handles scale-free graphs
- ▶ Handles large graphs
- ▶ Reduces computation not communication

# Partitioning Strategy

Placing high-degree vertices in one type of processor  
and low-degree ones in the other type.

- ▶ Partitions have different degrees of parallelism
- ▶ Partitions are more homogenous in terms of vertex connectivity
- ▶ Low cost in terms of computational and space complexity



# Three Different Partitioning Strategies

- ▶ **HIGH**: highest-degree vertices to CPU, lowest-degree vertices to GPU
- ▶ **LOW**: lowest-degree vertices to CPU, highest-degree vertices to GPU
- ▶ **RAND**: partitions the graph randomly

# Evaluation


- ▶ Breadth-First Search (BFS): based on the level-synchronous one by Hong et al.
- ▶ PageRank: algorithm used by Google Search to rank websites in their search engine results by assigning weights to hyperlinks

# Evaluation

- ▶ Scale-25 and scale-28 power-law degree distribution graphs
- ▶ Scale-25 and scale-28 uniform degree distribution graphs
- ▶ Scale-25 to scale-29 graphs different hardware configurations

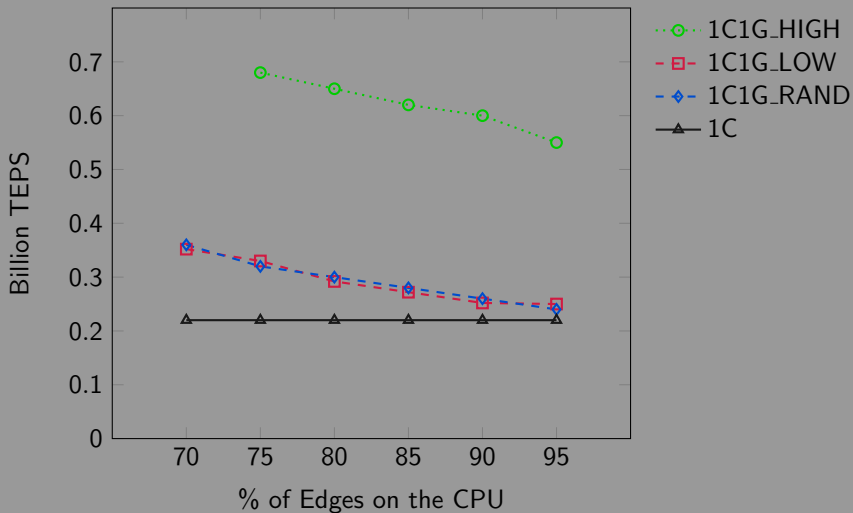
# Evaluation

BFS for scale-25 scale-free graph




# Evaluation

BFS for scale-28 scale-free graph




# Evaluation




# Evaluation

PageRank for scale-28 scale-free graph



# Evaluation

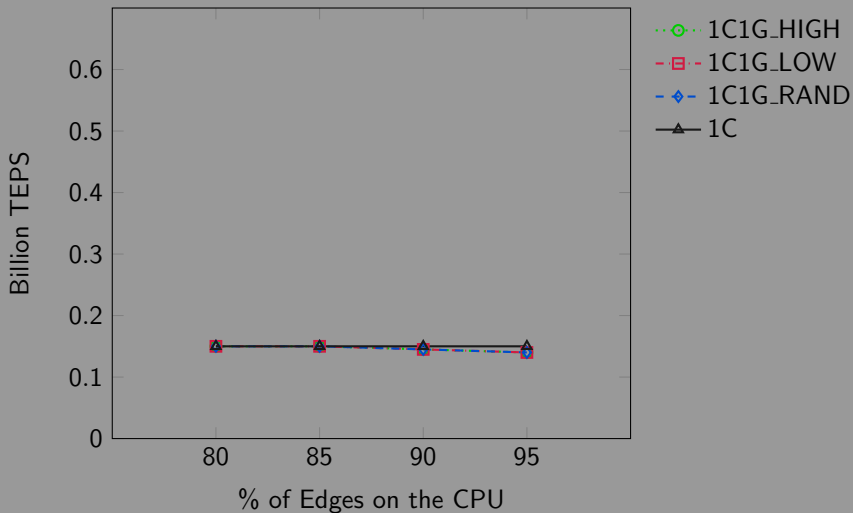
BFS for scale-25 uniform graph





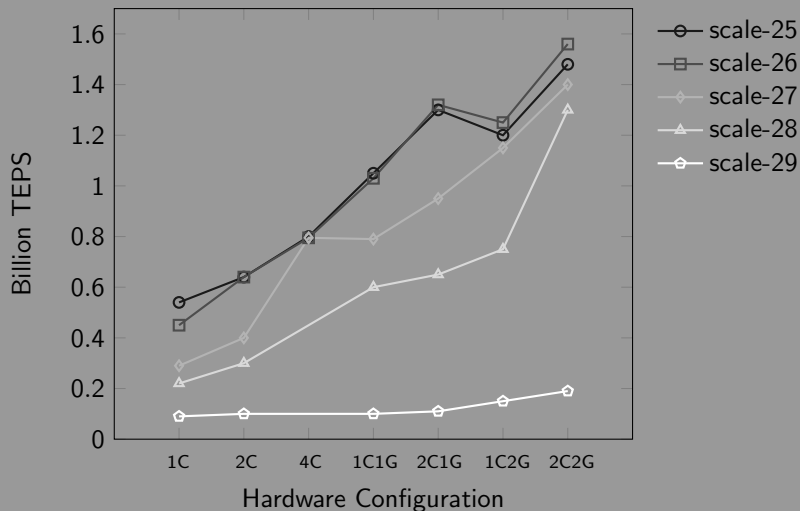
# Evaluation

BFS for scale-28 uniform graph



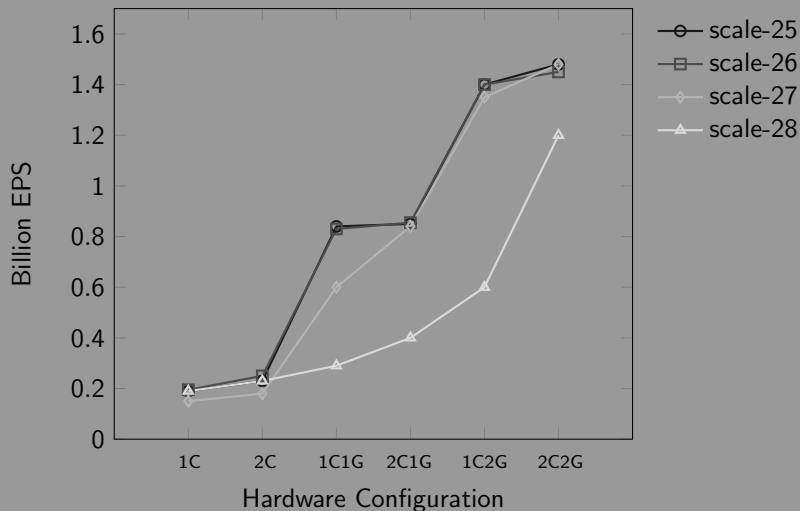
# Evaluation

BFS for scale-25 to scale-29 scale-free graphs



# Evaluation

PageRank for scale-25 to scale-28 scale-free graphs



# Summary

- ▶ Cache-sensitive algorithms:
  - ▶ Large graphs: HIGH
  - ▶ Small graphs: LOW
- ▶ More compute-intensive algorithms:
  - ▶ Large graphs: LOW
  - ▶ Small graphs: HIGH
- ▶ Hybrid-system makes sense
- ▶ Graph topology is important

# Personal Review

- ▶ Simple partition strategy
- ▶ Show in great detail how it improves performance
- ▶ Focus is on BFS
- ▶ Code is available

# Questions?